

# Day1

## Table of Contents

- [1. Agenda](#)
- [2. MPI Communicators](#)
- [3. Types of MPI Communications](#)
  - [3.1. Point-to-Point Communication:](#)
  - [3.2. Collective Communication:](#)
- [4. Sample MPI code](#)
- [5. Scripts](#)
  - [5.1. compile script](#)
  - [5.2. run script](#)
- [6. Hello World in C](#)
- [7. Hello World just by using MPI](#)
- [8. Hello World with MPI routines](#)
- [9. Hello from only even rank of processes](#)
- [10. MPI Initialization: MPI\\_Init vs. MPI\\_Initthread](#)
  - [10.1. Levels of Thread Support](#)
  - [10.2. MPI\\_Init Example](#)
  - [10.3. Compilation and Execution \(MPI\\_Init.\)](#)
  - [10.4. MPI\\_Initthread Example](#)
  - [10.5. Compilation and Execution \(MPI\\_Initthread.\)](#)
  - [10.6. Summary](#)
- [11. Point-to-point communication](#)
  - [11.1. Sending array to process 1](#)

## 1. Agenda

- Recap
- Communicators
- MPI Communications

- MPI Programs
- Point to point communications

## 2. MPI Communicators

Communicators in MPI define a group of processes that can communicate with each other. The default communicator is ``MPI_COMM_WORLD``, which includes all the processes. Custom communicators can be created to define subgroups of processes for specific communication patterns.

## 3. Types of MPI Communications

MPI offers various communication mechanisms to facilitate different types of data exchanges between processes:

### 3.1. Point-to-Point Communication:

- **Blocking:** The sending and receiving operations wait until the message is delivered (e.g., ``MPI_Send``, ``MPI_Recv``).
- **Non-Blocking:** The operations return immediately, allowing computation and communication to overlap (e.g., ``MPI_Isend``, ``MPI_Irecv``).

### 3.2. Collective Communication:

These operations involve a group of processes and include:

- **Broadcast:** Send data from one process to all other processes (``MPI_Bcast``).
- **Scatter:** Distribute distinct chunks of data from one process to all processes (``MPI_Scatter``).
- **Gather:** Collect chunks of data from all processes to one process (``MPI_Gather``).
- **All-to-All:** Every process sends and receives distinct chunks of data (``MPI_Alltoall``).

Collectives can also include operations like reductions (``MPI_Reduce``, ``MPI_Allreduce``) which perform computations on data from all processes and distribute the result.

## 4. Sample MPI code

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char **argv) {
    MPI_Init(&argc, &argv);

    int size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    int rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    printf("Hello from process %d of %d\n", rank, size);

    MPI_Finalize();
    return 0;
}
```

This example demonstrates a simple MPI program where each process prints its rank and the total number of processes.

## 5. Scripts

### 5.1. compile script

```
#!/bin/sh

#source /opt/ohpc/pub/apps/spack/share/spack/setup-env.sh
#spack load gcc/5i5y5cb
#spack load openmpi/c7kvqyq
source ~/git/spack/share/spack/setup-env.sh
spack load openmpi

inputFile=$1
outputFile="${1%.*}.out"      # extract the name of the file without extension and adding extension .out
#cmd=`mpicc $inputFile -o $outputFile`
cmd="mpicc $inputFile -o $outputFile"    # running code using MPI
```

```

echo "-----"
echo "Command executed: $cmd"
echo "-----"
$cmd

echo "Compilation successful. Check at $outputFile"
echo "-----"

```

## 5.2. run script

```

#!/bin/sh

#source /opt/ohpc/pub/apps/spack/share/spack/setup-env.sh
#spack load gcc/5i5y5cbc
source ~/git/spack/share/spack/setup-env.sh
spack load openmpi

cmd="mpirun -np $2 $1"
echo "-----"
echo "Command executed: $cmd"
echo "-----"
echo "#####"
echo "#####          OUTPUT          #####"
echo "#####"
echo
mpirun -np $2 $1
echo
echo "#####"
echo "#####          DONE          #####"
echo "#####"

```

## 6. Hello World in C

```

#include<stdio.h>
int main(){
    printf("Hello World\n");
    return 0;
}

```

```
Hello World
```

## 7. Hello World just by using MPI

```
#include<stdio.h>
int main(){
    printf("Hello World\n");
    return 0;
}
```

```
Hello World
```

```
bash compile.sh hello_mpi.c
```

```
-----
Command executed: mpicc hello_mpi.c -o hello_mpi.out
-----
Compilation successful. Check at hello_mpi.out
-----
```

```
bash run.sh ./hello_mpi.out 6
```

```
-----
Command executed: mpirun -np 6 ./hello_mpi.out
-----
#####
#####          OUTPUT          #####
#####
#####

Hello World
Hello World
Hello World
Hello World
Hello World
```

```
Hello World
```

```
#####  
#####          DONE          #####  
#####
```

## 8. Hello World with MPI routines

```
#include<stdio.h>  
#include<mpi.h>      // for using mpi functions  
int main(){  
    int size, rank;  
    MPI_Init(NULL, NULL);  
    MPI_Comm_size(MPI_COMM_WORLD, &size);  
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
    printf("Hello from process %d\n", rank);  
    MPI_Finalize();  
    return 0;  
}
```

```
bash compile.sh hello_mpi_processes.c
```

```
-----  
Command executed: mpicc hello_mpi_processes.c -o hello_mpi_processes.out  
-----  
Compilation successful. Check at hello_mpi_processes.out  
-----
```

```
bash run.sh ./hello_mpi_processes.out 8
```

```
-----  
Command executed: mpirun -np 8 ./hello_mpi_processes.out  
-----
```

```
#####  
#####          OUTPUT          #####  
#####
```

```
Hello from process 5
```

```
Hello from process 4
Hello from process 6
Hello from process 7
Hello from process 2
Hello from process 1
Hello from process 3
Hello from process 0
```

```
#####
#####          DONE          #####
#####
```

## 9. Hello from only even rank of processes

```
#include<stdio.h>
#include<mpi.h>      // for using mpi functions
int main(){
    int size, rank;
    MPI_Init(NULL, NULL);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if(rank % 2 == 0)
        printf("Hello from process %d\n", rank);
    MPI_Finalize();
    return 0;
}
```

```
bash compile.sh hello_mpi_even_processes.c
```

```
-----
Command executed: mpicc hello_mpi_even_processes.c -o hello_mpi_even_processes.out
-----
Compilation successful. Check at hello_mpi_even_processes.out
-----
```

```
bash run.sh ./hello_mpi_even_processes.out 10
```

```
-----
```

```
Command executed: mpirun -np 10 ./hello_mpi_even_processes.out
```

```
-----  
#####  
#####          OUTPUT          #####  
#####  
  
Hello from process 0  
Hello from process 2  
Hello from process 6  
Hello from process 8  
Hello from process 4  
  
#####  
#####          DONE          #####  
#####
```

## 10. MPI Initialization: `MPI_Init` vs. `MPI_Initthread`

MPI provides two main functions to initialize the MPI environment: ``MPI_Init`` and ``MPI_Initthread``. The primary difference is that ``MPI_Initthread`` allows you to specify the desired level of thread support.

### 10.1. Levels of Thread Support

- ``MPI_THREAD_SINGLE``: Only one thread will execute.
- ``MPI_THREAD_FUNNELED``: The process may be multi-threaded, but only the main thread will make MPI calls.
- ``MPI_THREAD_SERIALIZED``: Multiple threads may make MPI calls, but only one at a time.
- ``MPI_THREAD_MULTIPLE``: Multiple threads may make MPI calls with no restrictions.

### 10.2. `MPI_Init` Example

This example uses ``MPI_Init`` to initialize the MPI environment.

```
#include <mpi.h>  
#include <stdio.h>  
  
int main(int argc, char** argv) {
```



```

// Initialize the MPI environment
MPI_Init(&argc, &argv);

// Get the number of processes
int world_size;
MPI_Comm_size(MPI_COMM_WORLD, &world_size);

// Get the rank of the process
int world_rank;
MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

// Print off a hello world message
printf("Hello world from processor %d out of %d processors\n", world_rank, world_size);

// Finalize the MPI environment.
MPI_Finalize();
return 0;
}

```

### 10.3. Compilation and Execution (MPI<sub>Init</sub>)

- Compile the program:

```
bash compile.sh mpi_init.c
```

```

-----
Command executed: mpicc mpi_init.c -o mpi_init.out
-----
Compilation successful. Check at mpi_init.out
-----

```

- Run the program:

```
bash run.sh ./mpi_init.out 6
```

```

-----
Command executed: mpirun -np 6 ./mpi_init.out
-----

```

```
#####
#####          OUTPUT          #####
#####

Hello world from processor 1 out of 6 processors
Hello world from processor 2 out of 6 processors
Hello world from processor 5 out of 6 processors
Hello world from processor 0 out of 6 processors
Hello world from processor 4 out of 6 processors
Hello world from processor 3 out of 6 processors

#####
#####          DONE          #####
#####
```

## 10.4. MPI\_Initthread Example

This example uses `MPI\_Initthread` to initialize the MPI environment with thread support.

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    int provided;

    // Initialize the MPI environment with thread support
    MPI_Init_thread(&argc, &argv, MPI_THREAD_MULTIPLE, &provided);

    // Check the level of thread support provided
    if (provided < MPI_THREAD_MULTIPLE) {
        printf("MPI does not provide required thread support\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    // Print off a hello world message
```

```

printf("Hello world from processor %d out of %d processors with thread support level %d\n", world_rank, world_size, provided);

// Finalize the MPI environment.
MPI_Finalize();
return 0;
}

```

## 10.5. Compilation and Execution (MPI\_Initthread)

- Compile the program:

```
bash compile.sh mpi_init_thread.c
```

```

-----
Command executed: mpicc mpi_init_thread.c -o mpi_init_thread.out
-----
Compilation successful. Check at mpi_init_thread.out
-----

```

- Run the program:

```
bash run.sh ./mpi_init_thread.out 5
```

```

-----
Command executed: mpirun -np 5 ./mpi_init_thread.out
-----
#####
#####          OUTPUT          #####
#####
Hello world from processor 0 out of 5 processors with thread support level 3
Hello world from processor 2 out of 5 processors with thread support level 3
Hello world from processor 3 out of 5 processors with thread support level 3
Hello world from processor 1 out of 5 processors with thread support level 3
Hello world from processor 4 out of 5 processors with thread support level 3

#####
#####          DONE          #####
#####

```

```
#####
```

## 10.6. Summary

- `MPI_Init` is used for standard MPI initialization without considering threading.
- `MPI_Initthread` allows the program to specify and check the level of thread support.
  - Important for applications that require multi-threading in conjunction with MPI.
  - Ensures that the required thread support is available.

## 11. Point-to-point communication

```
#include "stdio.h"
#include "mpi.h"

int main(int argc, char **argv)
{
    int myid, size;
    int myval;
    MPI_Status status;

    //Initialize MPI environment
    MPI_Init(&argc,&argv);

    //Get total number of processes
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    //Get my unique ID among all processes
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);

    // Process with ID exactly equal to 0
    if(myid==0){
        //Initialize data to be sent
        myval = 100;
        //Print the data to be sent
        printf("\nmyid: %d \t myval = %d", myid, myval);
        //Send data
        MPI_Send(&myval, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
        printf("\nmyid: %d \t Data sent.\n", myid);
    }
    else if(myid==1){ // Process with ID exactly equal to 1
```

```

        //Initialize receive array to some other data
        myval = 200;
        MPI_Recv(&myval, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("\nmyid: %d \t Data received.", myid);
        //Print received data
        printf("\nmyid: %d \t myval = %d", myid, myval);
        printf("\n\nProgram exit!\n");
    }

    //End MPI environment
    MPI_Finalize();
}

```

```
bash compile.sh p2p_mpi.c
```

```

-----
Command executed: mpicc p2p_mpi.c -o p2p_mpi.out
-----
Compilation successful. Check at p2p_mpi.out
-----

```

```
bash run.sh ./p2p_mpi.out 2
```

```

-----
Command executed: mpirun -np 2 ./p2p_mpi.out
-----

```

```

#####
#####                                OUTPUT                                #####
#####
#####

```

```

myid: 0          myval = 100
myid: 0          Data sent.

myid: 1          Data received.
myid: 1          myval = 100

```

```
Program exit!
```

```

#####
#####                                DONE                                #####
#####

```

```
#####
```

## 11.1. Sending array to process 1

```
#include "stdio.h"
#include "mpi.h"
#define N 100

int main()
{
    int myid, size;
    int myval;

    int arr[N];
    //Initialize MPI environment
    MPI_Init(NULL, NULL);

    //Get total number of processes
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    //Get my unique ID among all processes
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);

    // Process with ID exactly equal to 0
    if(myid==0){
        //Initialize data to be sent
        for(int i = 0; i < N; i++) arr[i] = i + 1;
        //Send data
        MPI_Send(arr, N, MPI_INT, 1, 0, MPI_COMM_WORLD);
        printf("\nmyid: %d \t Data sent.\n", myid);
    }
    else if(myid==1){ // Process with ID exactly equal to 1
        //Initialize receive array to some other data
        MPI_Recv(arr, N, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("\nmyid: %d \t Data received.\n", myid);
        //Print received data
        for(int i = 0; i < N; i++)
            printf("%d ", arr[i]);
    }

    //End MPI environment
    MPI_Finalize();
}
```

```
bash compile.sh p2p_mpi_array.c
```

```
-----  
Command executed: mpicc p2p_mpi_array.c -o p2p_mpi_array.out  
-----
```

```
Compilation successful. Check at p2p_mpi_array.out  
-----
```

```
bash run.sh ./p2p_mpi_array.out 2
```

```
-----  
Command executed: mpirun -np 2 ./p2p_mpi_array.out  
-----
```

```
#####  
#####                OUTPUT                #####  
#####
```

```
myid: 1          Data received.
```

```
myid: 0          Data sent.
```

```
1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48  
#####  
#####                DONE                #####  
#####
```

Author: Abhishek Raj

Created: 2024-07-01 Mon 16:59